Lecture 6: Hashing II: Table Doubling, Rolling Hash, Karp-Rabin

Lecture Overview

- Table Resizing

- Amortization

- DNA Comparison, Karp-Rabin Rolling Hash

Readings

CLRS Chapter 17 and 32.2.

**Recall:**

- **Hashing with Chaining:**

Figure 1: Chaining in a Hash Table

- **Simple Uniform Hashing Assumption (Silently used in all of this lecture)**

Each key is equally likely to be hashed to any slot of table, independent of where other keys are hashed.

⇒ average ♯ keys per slot is $\alpha = n/m$

⇒ expected time to search, insert, delete $= O(1 + \alpha)$.

Caution: The above bound assumes that the application of the hash function $h(\cdot)$ takes $O(1)$ time. Sometimes this is not the case, e.g. if the keys are strings. Then the keys need to be processed into numbers and then hashed to $\{0, 1, \ldots, m-1\}$. Then the above bound is scaled by the time needed for applying h.

- **Good Hash Functions:**

  – Division Method: $h(k) = k \bmod m$

    Good Practice: m is a prime number & not close to a power of 2 or 10.

  – Multiplication Method: $h(k) = [(a \cdot k) \bmod 2^w] \gg (w - r)$, where

    - $\gg$ denotes the "shift right" operator,

    - 2^r is the table size $(= m)$,

    - w the bit-length of the machine words,

    - and a is chosen to be an odd integer between $2^{(w-1)}$ and 2^w.

   *Good Practice:* a not too close to $2^{(w-1)}$ or 2^w.

Figure 2: Multiplication Method

How Large should Table be?

- want $m = \Theta(n)$ at all times

- **Why?**
  m too small $\implies$ slow (recall operations take expected time $\Theta(1 + \alpha)$, where $\alpha = \frac{n}{m}$);
  m too big $\implies$ wasteful

Challenge: Don't know how large n will get at creation.

Idea: Start small (constant) and grow (or shrink) as necessary.

Table Resizing with Rehashing:

To change m build new hash table from scratch:

Allocate table of size m;
For each item in old table: $\implies \Theta(n + m)$ time $= \Theta(n)$, if $m = \Theta(n)$
  insert into new table

How fast to grow?

When n reaches m, say

- **$m\ += 1$?**
  $\implies$ rebuild every step
  $\implies$ n inserts cost $\Theta(1 + 2 + \cdots + n) = \Theta(n^2)$

- **$m\ *= 2$?** $m = \Theta(n)$ still
  $\implies$ rebuild at insertion 2^i, pay 2^{i+1} (see Figure 3)
  $\implies$ n inserts cost $\Theta(1 + 2 + 4 + 8 + \cdots + n)$ where n is really the next power of 2
  $= \Theta(n)$

- a few inserts cost linear time, but $\Theta(1)$ "on average".

Amortized Analysis

This is a common technique in data structures - like paying rent: \$ 1500/month $\approx$ \$ 50/day

- if a sequence of n operations has total cost $\leq n \cdot T(n)$, then each operation has *amortized cost $T(n)$*

- "$T(n)$ amortized" roughly means $T(n)$ "on average", but averaged over all ops.

- e.g. inserting into a hash table (with doubling) takes $O(1)$ amortized time.

Figure 3: Resizing by Doubling. Costly Inserts: Item # 1, 2, 4, 8, 16, ... Resizing Penalty is constant on average.

Back to Hashing:

Maintain $m = \Theta(n)$ so also support search in $O(1)$ expected time assuming simple uniform hashing.

Delete:

Also $O(1)$ expected time

- space can get big with respect to n e.g. $n \times$ insert, $n \times$ delete

- <u>solution</u>: when n decreases to $m/4$, shrink table to $m/2$
       $\implies O(1)$ amortized cost for both insert and delete —
          —analysis is trickier; (see CLRS 17.4).

Rolling Hash: Human vs Chimp

Given two strings S and T, find the longest common substring of two strings.

Naive algorithm: $\Theta(n^4)$.

Naive + binary search: $\Theta(n^3 \log n)$.

Winner algorithm from last lecture runs in time $\Theta(n^2 \log n)$, using hash tables:

For all possible lengths ℓ:

- (Step 1) Insert all substrings s of S of length ℓ into a hash table using some hash function h;

- (Step 2) For all substrings t of T of length ℓ, check if position $h(t)$ of the dictionary is occupied. If yes (say by substring s of string S), compare s and t. If strings agree return s and t and exit.

Analysis:

Outer Loop: using binary search on the length of the longest common substring, only $O(\log n)$ iterations are needed!

For every ℓ:

  Step 1:

- there are $n - \ell + 1$ substrings s of S of length ℓ;

- need to convert each of them into an integer; **how?**
  think of $s := S[i \ : \ i+\ell]$ as a multi-digit number base b, where b is larger than the alphabet size

$$s \mapsto S[i] \cdot b^{\ell-1} + S[i+1] \cdot b^{\ell-2} + \ldots + S[i+\ell-1]$$

- use hash function $h(s) = s \bmod m$

- **How to compute $h(s)$ without writing down the above expression?**
  Mod arithmetic magic!

  Claim: For all integers a, b:

  $$a + b \bmod m = ((a \bmod m) + (b \bmod m)) \bmod m$$

  $$a \cdot b \bmod m = ((a \bmod m) \cdot (b \bmod m)) \bmod m$$

- hence, computation of hash takes time $O(\ell)$ for each substring s.
  $\implies$ total time for Step 1: $O(\ell \cdot (n - \ell + 1)) = O(n^2)$.

Step 2:

- there are $n - \ell + 1$ substrings t of T of length ℓ;

- every hash operation takes $O(\ell)$, plus another potential $O(\ell)$ for comparing strings if hashes match.

- $\implies$ total time for Step 1: $O(\ell \cdot (n - \ell + 1)) = O(n^2)$.

OVERALL time is $O(n^2 \log n)$.

OUR GOAL: Drop time down to $O(n^2)$ initialization time $+ O(n \log n)$ execution time. By making both Step 1 and Step 2 take $O(n)$.

Idea: use $h1 := h(S[i \; : \; i + \ell])$ to compute $h2 := h(S[i + 1 \; : \; i + \ell + 1])$ in $O(1)$ time.

How?

Go from

$$h1 = S[i] \cdot b^{\ell-1} + S[i+1] \cdot b^{\ell-2} + \ldots + S[i + \ell - 1] \bmod m$$

to

$$h2 = S[i+1] \cdot b^{\ell-1} + S[i+2] \cdot b^{\ell-2} + \ldots + S[i + \ell] \bmod m.$$

without recomputing $h2$ from scratch.

Magic Again — Called Rolling Hash, and introduced by Karp-Rabin:

$$h2 = \Big(S[i+1] \cdot b^{\ell-1} + S[i+2] \cdot b^{\ell-2} + \ldots + S[i+\ell] \Big) \bmod m =$$
$$= \Big(\big(S[i] \cdot b^{\ell-1} + S[i+1] \cdot b^{\ell-2} + \ldots + S[i + \ell - 1] \big) b + S[i+\ell] - S[i] \cdot b^{\ell} \Big) \bmod m =$$
$$= \Big(h1 \cdot b + S[i+\ell] - S[i] \cdot (b^{\ell} \bmod m) \Big) \bmod m.$$

Now Step 1 takes time $O(n)$ overall. **What about Step 2?** Also, time $O(n)$, except if there are too many *spurious* matches of hash values, which do not actually result in matching substrings: Each comparison takes $O(\ell)$ and if many unsuccessful comparisons are made this could increase the time for Step 2 to $O(n\ell)$.

But: For each substring of T the probability of a false-positive hash-value match is $< \frac{n}{m}$ (under the simple uniform hashing assumption). Hence, choosing $m = O(n^2)$, the expected time to process a substring t that does not match a substring of S is $< \frac{1}{n} \cdot O(\ell) = O(1)$ ($1/n$ is the probability that the hash value of t collides with the hash value of a substring of S and $O(\ell)$ is the time to carry out the string comparison).

Hence, $O(n^2)$ time to allocate memory for a table of size $m = O(n^2)$
$+ O(n \log n)$ execution time.

NOTE: We can get away with a hash table of size $\Theta(n)$, rather than $\Theta(n^2)$, while avoiding a large comparison cost from *spurious* hash-value matches with a further trick: if a substring s occupies position $h(s)$ of the hash table, we store at that position both s and a *signature* of s, produced by taking the multi-digit number corresponding to the string mod p, where p is a prime number larger than ℓ (and different than m). Now if position $h(t)$ of the hash table contains a substring s then we first compare the signatures of s and t, and only if these agree (probability $< 1/p$ for spurious matches under the simple uniform hashing assumption for the signature hash function), we compare the strings s and t. In this case, the expected time to process a spurious match is $< 1/p \cdot O(\ell) = O(1)$, even if the hash table has size $O(n)$. To avoid incurring $O(\ell)$ cost for each signature computation, we do rolling hashing for the signatures as well :-). In this computation, we implicitly assumed that the values produced by the hash function h and the signature function are independent, and that both functions satisfy the simple uniform hashing assumption.